

Assignment 2

Learning objectives

1. Implement a correct, robust AVL (self-balancing binary search) tree from scratch.
2. Understand and implement single (LL, RR) and double (LR, RL) rotations.
3. Apply AVL trees to a practical problem and measure empirical performance.
4. Compare AVL behavior against an unbalanced BST (and language built-ins when available).

Overview

In this assignment you will implement an AVL tree that supports insertion, deletion, search, and traversal. You will then use your AVL implementation in a small application: a **word-frequency manager** that stores words (strings) and their counts. Finally, you will run experiments comparing AVL against a simple unbalanced BST implementation and report the results.

Requirements (mandatory)

1. **AVL Implementation**
 - o Implement an AVL tree class `AVLTree` with nodes containing a key and associated value (for the word-frequency manager the key is a string and value is an integer). Your implementation must include:
 - `insert(key, value)` — if the key exists, update the value (or allow an increment operation; choose one and document it).
 - `delete(key)` — remove a key if present.
 - `search(key)` — return associated value or `null/None/optional` if not present.
 - `in_order_traversal()` — return a list/array of (key, value) pairs in ascending key order.
 - Access to `height(node)` and `balance_factor(node)` (internally) and correct rotations: LL, RR, LR, RL.
 - o The tree must maintain AVL balance invariants after every insertion and deletion.
2. **Unbalanced BST Implementation**
 - o Implement a plain Binary Search Tree `BST` with the same public interface as `AVLTree` (`insert`, `delete`, `search`, `traversal`) for fair comparison. You may reuse parts of code but must not include balancing logic.
3. **Application: Word-Frequency Manager**
 - o Read a text file (plain `.txt`) and insert all words into both data structures (AVL and BST) in the same order.
 - o Normalize words by: lowercasing and removing punctuation (define your normalization clearly).
 - o For each insertion, if the word already exists, increment its frequency count.

- o Provide functionality to:
 - Query frequency of a given word.
 - Output the top- k frequent words (you may compute this by traversing and sorting the (word, count) list — justify your complexity in the report).
 - Persist the in-order traversal to a file out.txt (one word count per line).
- 4. **Performance Experiments**
 - o Design and run experiments to compare AVL vs BST on the following metrics:
 - **Average time per insertion** (measured in milliseconds or microseconds). Use at least three different input sizes (e.g., 10k, 50k, 200k words — adjust to machine capability).
 - **Average tree height** after all insertions.
 - **Time for search queries**: pick a set of random words (present and absent) and report mean search time.
 - o Run each experiment multiple times and report averaged results. Present results as a table and (optionally) simple plots.
- 5. **Report** (Maximum 2 pages)
 - o Present results (tables/plots) and analyze them: explain why AVL performed better/worse in different metrics.
 - o Include complexity analysis (worst-case/average-case) for insertion, deletion, search, and obtaining top- k .

Deliverables

1. Source code (single zip or repository) including:
 - o AVLTree implementation.
 - o BST implementation.
 - o Application scripts (word-frequency manager, experiment runner).
 - o Unit tests.
2. report.pdf